

Statistical Learning Lab: Assignment-04

Cross Validation and Bootstrapping

Name: Semanti Ghosh

Roll No: 22IM10036

1. Loading the data and printing the first few lines

Code snippet

```
1 #getting to the current working directory
2 setwd("C:/Study/Semester_6/Statistical_Learning_Lab")
3 getwd()
4 |
5 #reading and printing the data
6 data <- read.csv("manufacturing.csv")
7 head(data)
```

Output

```
> #getting to the current working directory
> setwd("C:/Study/Semester_6/Statistical_Learning_Lab")
> getwd()
[1] "C:/Study/Semester_6/Statistical_Learning_Lab"
>
> #reading and printing the data
> data <- read.csv("manufacturing.csv")
> head(data)
  Temperature...C. Pressure..kPa. Temperature.x.Pressure Material.Fusion.Metric
1      209.7627      8.050855      1688.769      44522.22
2      243.0379     15.812068      3842.931      63020.76
3      220.5527      7.843130      1729.823      49125.95
4      208.9766     23.786089      4970.737      57128.88
5      184.7310     15.797812      2918.345      38068.20
6      229.1788      8.498306      1947.632      53136.69
  Material.Transformation.Metric Quality.Rating
1           9229576           99.99997
2          14355367           99.98570
3          10728389           99.99976
4           9125702           99.99997
5           6303792          100.00000
6          12037072           99.99879
>
```

2. Fitting polynomials on temperature from 1 to 5 degrees and then performing LOOCV and k-Fold CV

Code Snippet

```
#varying the degree of polynomial on temperature
library(ggplot2)
models <- list()
degrees <- 1:5

for (d in degrees) {
  formula <- as.formula(paste("Quality.Rating ~ poly(Temperature...C.,", d, ", raw=TRUE)"))
  models[[d]] <- lm(formula, data = data)
}

library(boot)

#initialization
cv_results <- data.frame(Degree=integer(), LOOCV=numeric(), k5=numeric(), k10=numeric())

#performing LOOCV and k-fold CV (doing it for different degrees)
for (d in degrees) {
  formula <- as.formula(paste("Quality.Rating ~ poly(Temperature...C.,", d, ", raw=TRUE)"))
  model <- glm(formula, data = data)

  loocv_error <- cv.glm(data, model)$delta[1]
  k5_error <- cv.glm(data, model, K=5)$delta[1]
  k10_error <- cv.glm(data, model, K=10)$delta[1]

  cv_results <- rbind(cv_results, data.frame(Degree=d, LOOCV=loocv_error, k5=k5_error, k10=k10_error))
}

#Displaying the results in a table
install.packages("pander")
library(pander)
pander(cv_results)
```

Output

```
>
> for (d in degrees) {
+   formula <- as.formula(paste("Quality.Rating ~ poly(Temperature...C.,", d, ", raw=TRUE)"))
+   models[[d]] <- lm(formula, data = data)
+ }
>
> library(boot)
>
> #initialization
> cv_results <- data.frame(Degree=integer(), LOOCV=numeric(), k5=numeric(), k10=numeric())
>
> #performing LOOCV and k-fold CV (doing it for different degrees)
> for (d in degrees) {
+   formula <- as.formula(paste("Quality.Rating ~ poly(Temperature...C.,", d, ", raw=TRUE)"))
+   model <- glm(formula, data = data)
+
+   loocv_error <- cv.glm(data, model)$delta[1]
+   k5_error <- cv.glm(data, model, K=5)$delta[1]
+   k10_error <- cv.glm(data, model, K=10)$delta[1]
+
+   cv_results <- rbind(cv_results, data.frame(Degree=d, LOOCV=loocv_error, k5=k5_error, k10=k10_error))
+ }
> library(pander)
> pander(cv_results)
```

```
-----
Degree  LOOCV    k5    k10
-----
1      133.1   133.5   133.1
2       91.19   91.09   91.09
3       54.24   54.2    54.15
4       28.79   28.69   28.76
5       13.53   13.56   13.49
-----
```


Plotting the results

Code snippet

```
# plotting the results
cv_results_long <- reshape2::melt(cv_results, id="Degree")

ggplot(cv_results_long, aes(x=Degree, y=value, color=variable)) +
  geom_line() + geom_point() +
  labs(title="Cross-Validation MSE for Different Polynomial Degrees",
       x="Polynomial Degree", y="MSE", color="CV Method") +
  theme_minimal()
```

Output



Analysis: From the above graph and results, we can say that when we fit a polynomial of degree 5, then the MSE is minimised. This is expected since the training and validation set error does decrease when the degree of the polynomial increases, since then the polynomial becomes more flexible. However, we should also keep in mind that if the degree of the polynomial is too high, then it will overfit the data, resulting in high variance. So, while the results suggest that 5 is the best choice, we might choose a lower degree polynomial (like maybe three or four degree polynomial). This will make the model more robust and ensure generalisation on unseen data.

3. Fitting linear polynomials for different variables

Code Snippet

```
#trying out for linear combinations of different variables(excluding interactions)
library(reshape2)

# Define all variable combinations
var_combinations <- list(
  "Temp" = c("Temperature...C."),
  "Temp-Press" = c("Temperature...C.", "Pressure..kPa."),
  "Temp-MatFus" = c("Temperature...C.", "Material.Fusion.Metric"),
  "Temp-MatTrans" = c("Temperature...C.", "Material.Transformation.Metric"),
  "Temp-Press-MatFus" = c("Temperature...C.", "Pressure..kPa.", "Material.Fusion.Metric"),
  "Temp-Press-MatTrans" = c("Temperature...C.", "Pressure..kPa.", "Material.Transformation.Metric"),
  "Temp-MatFus-MatTrans" = c("Temperature...C.", "Material.Fusion.Metric", "Material.Transformation.Metric"),
  "Temp-Press-MatFus-MatTrans" = c("Temperature...C.", "Pressure..kPa.", "Material.Fusion.Metric", "Material.Transformation.Metric")
)

cv_results <- data.frame(Model=character(), LOOCV=numeric(), k5=numeric(), k10=numeric())

for (model_name in names(var_combinations)) {
  formula <- as.formula(paste("Quality.Rating ~", paste(var_combinations[[model_name]], collapse = " + ")))
  model <- glm(formula, data = data)

  loocv_error <- cv.glm(data, model)$delta[1]
  k5_error <- cv.glm(data, model, K=5)$delta[1]
  k10_error <- cv.glm(data, model, K=10)$delta[1]

  cv_results <- rbind(cv_results, data.frame(Model=model_name, LOOCV=loocv_error, k5=k5_error, k10=k10_error))
}

#printing the results
print(cv_results)
```

Output

```
#trying out for linear combinations of different variables(excluding interactions)
library(ggplot2)
library(reshape2)

# Define all variable combinations
var_combinations <- list(
  "Temp" = c("Temperature...C."),
  "Temp-Press" = c("Temperature...C.", "Pressure..kPa."),
  "Temp-MatFus" = c("Temperature...C.", "Material.Fusion.Metric"),
  "Temp-MatTrans" = c("Temperature...C.", "Material.Transformation.Metric"),
  "Temp-Press-MatFus" = c("Temperature...C.", "Pressure..kPa.", "Material.Fusion.Metric"),
  "Temp-Press-MatTrans" = c("Temperature...C.", "Pressure..kPa.", "Material.Transformation.Metric"),
  "Temp-MatFus-MatTrans" = c("Temperature...C.", "Material.Fusion.Metric", "Material.Transformation.Metric"),
  "Temp-Press-MatFus-MatTrans" = c("Temperature...C.", "Pressure..kPa.", "Material.Fusion.Metric", "Material.Transformation.Metric")
)

cv_results <- data.frame(Model=character(), LOOCV=numeric(), k5=numeric(), k10=numeric())

for (model_name in names(var_combinations)) {
  formula <- as.formula(paste("Quality.Rating ~", paste(var_combinations[[model_name]], collapse = " + ")))
  model <- glm(formula, data = data)

  loocv_error <- cv.glm(data, model)$delta[1]
  k5_error <- cv.glm(data, model, K=5)$delta[1]
  k10_error <- cv.glm(data, model, K=10)$delta[1]

  cv_results <- rbind(cv_results, data.frame(Model=model_name, LOOCV=loocv_error, k5=k5_error, k10=k10_error))
}

# Print the CV results
print(cv_results)
```

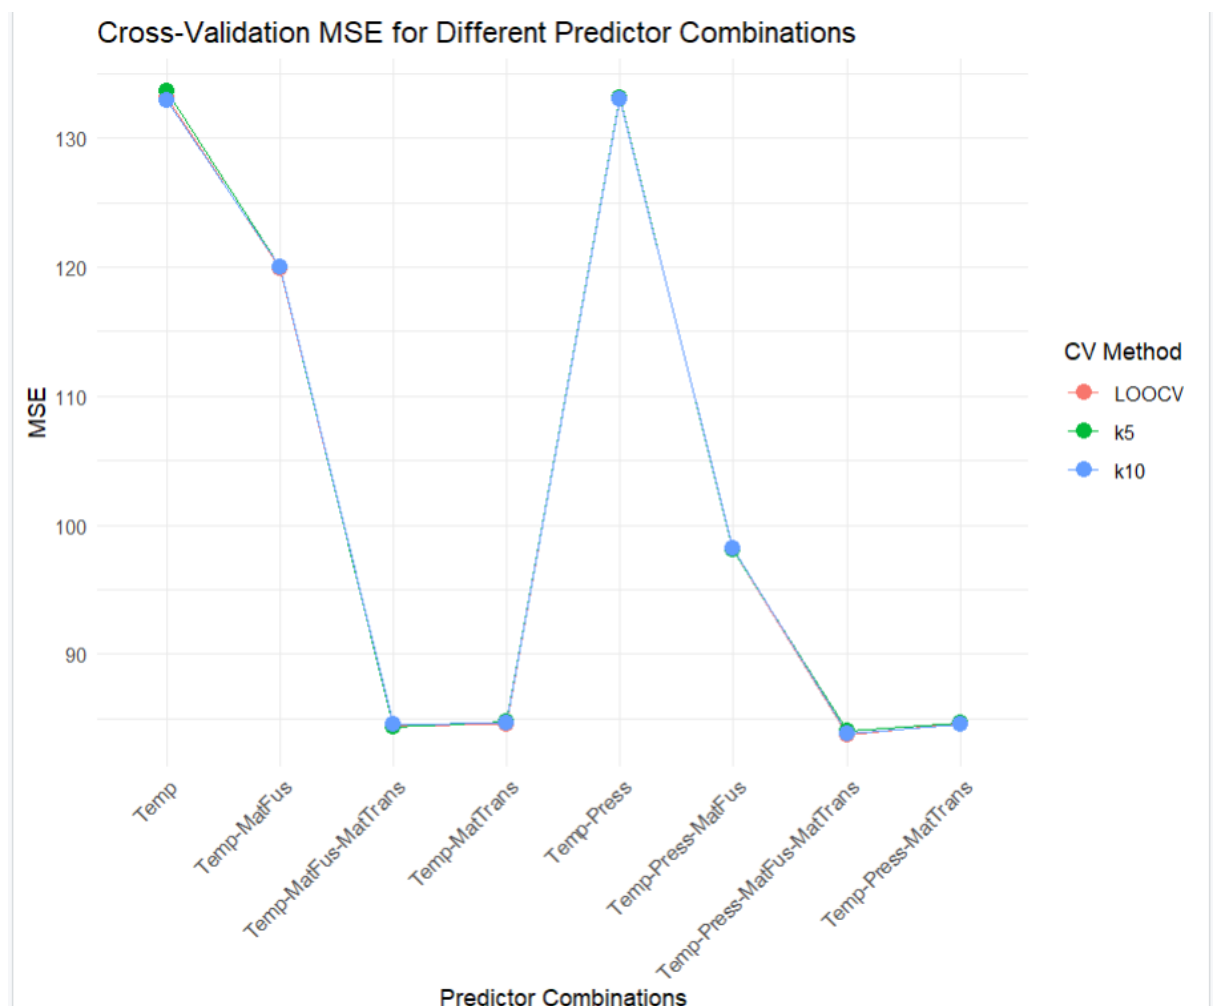
Model	LOOCV	k5	k10
Temp	133.07880	133.64967	132.91792
Temp-Press	133.14410	133.09729	133.04600
Temp-MatFus	119.93438	119.97232	119.96603
Temp-MatTrans	84.59484	84.80396	84.69679
Temp-Press-MatFus	98.12928	98.06263	98.20927
Temp-Press-MatTrans	84.63338	84.62915	84.53809
Temp-MatFus-MatTrans	84.50945	84.31112	84.56092
Temp-Press-MatFus-MatTrans	83.76925	84.04034	83.83754

Plotting the results

Code Snippet

```
cv_results_long <- melt(cv_results, id="Model")
|
ggplot(cv_results_long, aes(x=Model, y=value, color=variable)) +
  geom_line(aes(group=variable)) +
  geom_point(size=3) +
  labs(title="Cross-Validation MSE for Different Predictor Combinations",
       x="Predictor Combinations", y="MSE", color="CV Method") +
  theme_minimal() +
  theme(axis.text.x = element_text(angle=45, hjust=1)) # Rotate labels for readability
```

Output



Analysis: From the results obtained, we can say that the error is minimised when all four parameters (Temperature, Pressure, Material Fusion Metric and Material

Transformation Metric) are taken into consideration. However, like in the previous case, this can lead to overfitting resulting in high variance. So, even if the error is slightly greater, it's safer to go for three parameters, like in Temperature, Material Fusion Metric, Material Transformation Metric. This will make the model more robust and ensure generalisation on unseen data.

4. Bootstrapping

Code snippet

```
#Bootstrapping

#generating the random variables that follow Gaussian distribution
set.seed(123)
data <- rnorm(50, mean = 50, sd = sqrt(2))

bootstrap_means <- numeric(100)
bootstrap_vars <- numeric(100)

#random sampling with replacement
for (i in 1:100) {
  sample_data <- sample(data, size = 20, replace = TRUE)
  bootstrap_means[i] <- mean(sample_data)
  bootstrap_vars[i] <- var(sample_data)
}

#estimating mean and variance from the
boot_mean_estimate <- mean(bootstrap_means)
boot_var_estimate <- mean(bootstrap_vars)

cat("Estimated Population Mean from Bootstrap Samples:", boot_mean_estimate, "\n")
cat("Estimated Population Variance from Bootstrap Samples:", boot_var_estimate, "\n")
```

Output

```
> #Bootstrapping
> 
> #generating the random variables that follow Gaussian distribution
> set.seed(123)
> data <- rnorm(50, mean = 50, sd = sqrt(2))
> 
> bootstrap_means <- numeric(100)
> bootstrap_vars <- numeric(100)
> 
> #random sampling with replacement
> for (i in 1:100) {
+   sample_data <- sample(data, size = 20, replace = TRUE)
+   bootstrap_means[i] <- mean(sample_data)
+   bootstrap_vars[i] <- var(sample_data)
+ }
> 
> #estimating mean and variance from the
> boot_mean_estimate <- mean(bootstrap_means)
> boot_var_estimate <- mean(bootstrap_vars)
> 
> cat("Estimated Population Mean from Bootstrap Samples:", boot_mean_estimate, "\n")
Estimated Population Mean from Bootstrap Samples: 50.05668
> cat("Estimated Population Variance from Bootstrap Samples:", boot_var_estimate, "\n")
Estimated Population Variance from Bootstrap Samples: 1.733931
```

Result

Estimated Population Mean = 50.05668

Estimated Population Variance = 1.733931

Plotting the graph of means and variances in bootstrapping for better visualisation

Code Snippet

```
#plotting the bootstrapping means and variances (just for my own understanding)

p1 <- ggplot(data.frame(Means = bootstrap_means), aes(x = Means)) +
  geom_histogram(binwidth = 0.1, fill = "blue", color = "black", alpha = 0.7) +
  labs(title = "Bootstrap Sampling Distribution of the Mean",
       x = "Sample Mean", y = "Frequency") +
  theme_minimal()

p2 <- ggplot(data.frame(Variances = bootstrap_vars), aes(x = Variances)) +
  geom_histogram(binwidth = 0.05, fill = "red", color = "black", alpha = 0.7) +
  labs(title = "Bootstrap Sampling Distribution of the Variance",
       x = "Sample Variance", y = "Frequency") +
  theme_minimal()

print(p1)
print(p2)
```

Output

